

# SQL

David Zollikofer

This short paper summarizes important SQL concepts. All the content and examples are from *Database Systems: The Complete Book* by Garcia-Molina, Ullman and Widom as well as ETH’s lecture slides by Prof. Ce Zhang.

## Part I SQL for Queries

### Simple SQL Queries

A basic SQL query has the form `SELECT <attr> FROM <rels> WHERE <cond>` where

- `<attr>` are the attributes of the relation we would like to have
- `<rels>` are the tuples (the relation) we are selecting from
- `<cond>` filters the tuples in the relation only keeping the ones satisfying the condition

A basic query might look as follows:

```
SELECT title, year, producer FROM
movies WHERE id = 10;
```

Conditions are specified with `<attr-name> = <value>` and can be chained together using `AND`, `OR`, `NOT`, `()`

### String Comparison

Strings are lexicographically ordered in SQL and can be compared to each other.

We can also use wildcards:

- `_` matches exactly one character: `WHERE title LIKE 'Ice ___'` would match any movie that has three characters following the space. (such as Ice age or Ice tea).
- `%` matches 0 or more characters: `WHERE title LIKE '%sor%s'` matches or, order, etc.

### Dates and Times

Are prefixed with a keyword `DATE`, `TIMESTAMP`, `TIME`.

- `DATE '1999-12-01'` is the first of December of 1999
- `TIME '15:00:03.25'` is 3 and a quarter minutes past three PM.
- `TIMESTAMP '1966-03-22 13:00:12-1:00'`, notice the `-1:00` meaning one hour behind GTM.

To extract information you can use `SELECT EXTRACT(<selector> FROM <date/time>)` for example `SELECT EXTRACT(MONTH FROM '2000-00-12')`

### NULL Values and Unknowns

When a value is not specified in the table it defaults to `NULL`. When we make comparisons and have a `NULL` involved the truth value will become an `UNKNOWN`. On logical operators in behaves similarly to a `FALSE`, though it is not `FALSE`. Arithmetic expressions on `NULL` result in `NULL`.

The database only returns tuples where the `WHERE` clause evaluates to true. Hence if we have an `UNKNOWN` it will not be returned.

### Ordering

After a `WHERE`, `GROUP BY`, `HAVING` clause we can have an `ORDER BY` clause which looks as follows:

`ORDER BY <attrs> <order>`

where `<attrs>` is the list of attributes we want to order by and `<order>` is either `ASC` or `DESC` with the default being `ASC`.

### Union, Intersection and Differences

Let `(Q1)`, `(Q2)` be two queries of the form `SELECT ... FROM ... WHERE ....` We now define:

- `(Q1) UNION (Q2)` returns records either in `(Q1)` or `(Q2)` or both.
- `(Q1) INTERSECT (Q2)` returns only records which are in `(Q1)` and `(Q2)`
- `(Q1) EXCEPT (Q2)` returns all records being only in `(Q1)` and not in `(Q2)`.

### Case Statements

We can modify a row dynamically using the case statement. It is especially useful to deal with `NULL` values.

```
SELECT carId, carName,
CASE
    WHEN carAge = 0 THEN 'new'
    WHEN carAge < 20 THEN 'worthless'
    ELSE 'old-timer'
END AS ageComment
FROM cars;
```

If no `CASE` matches the statement will return `NULL`.

### Joins

### Cartesian Product / Cross Join

```
SELECT * FROM
Cars c CROSS JOIN Homeowners h
WHERE c.owner = h.id
```

Alternatively there also exist the traditional version with the implicit join:

```
SELECT * FROM
Cars c, Homeowners h
WHERE c.owner = h.id
```

This query build the cross product between all cards and all homeowners and keeps the records where the car owner identifier is the homeowner id.

**Note:** Sometimes one also sees `Cars as c`, `Homeowners as h` the keyword `as` is completely optional.

When `h` and `c` both have an attribute `id` we have to reference it as either `c.id` or `h.id` since otherwise there would be ambiguity.

### Normal (inner) Join

Combines tuples from both relations satisfying join condition.

```
SELECT * FROM
Cars c JOIN Pets p ON c.owner = p.owner
```

or equivalently if both have a column `owner` one would use `as` a join condition:

```
SELECT * FROM
Cars c JOIN Pets p USING(owner)
```

### Natural Joins

Combines tuples from both relations where all columns with the same name match. In the example if they both had a column `owner`, this column would have to match.

```
SELECT * FROM
Cars c NATURAL JOIN Pets p
```

### Outer Joins

They can be natural or normal. The key difference is that an outer join also includes non matching records and sets the missing attributes to `NULL`.

```
Cars c FULL JOIN Pets p using(owner)
```

The statement above combines all pet and car combinations as a normal join but also includes all tuples with no match. (all cars not sharing an owner with a cat and vice versa).

### Left Outer Join



```
Cars c LEFT OUTER JOIN Pets p using(owner)
```

The statement above includes all pet and car combinations as a normal join but also includes all cars not sharing an owner with a cat. But the cats not sharing an owner with a car are excluded. Analogously there exists a **Right Outer Join**.

### Subqueries

There are three places where SQL subqueries can occur:

- Returning a single value to be used in a WHERE clause.
- Returning relations to be used in a WHERE clause.
- In a FROM clause followed by a variable to represent that subquery.

### Single Value Returning

```
SELECT * FROM CarOwners co WHERE co.id = (SELECT owner
FROM Cars WHERE numberPlate = 123 LIMIT 1)
```

Notice that the LIMIT 1 is important. If multiple or no tuples were returned an error would be raised.

### Conditions with Relations

There are four types:

- $s > ALL\ R$ , is true if s is larger than all values in R.
- $s > ANY\ R$ , is true if is larger than at least one value in R.
- EXISTS R, is true if R is non-empty.
- $s\ IN\ R$ , is true if the tuple s is equal to a tuple in R. (analogously there is NOT IN)

As an example: (notice that the parentheses around owner, insurance are essential to mark it as a tuple). Here we find the car owners who have insurance for their car.

```
SELECT owner FROM CarOwners, Insurance WHERE (owner,
insurance) IN (SELECT owner, insurance FROM
OwnerInsurance)
```

### Correlated Subquery

We use the tuple we are selecting in a subquery:

```
SELECT * FROM Cars c
WHERE productionYear <= ALL (
    SELECT year FROM Cars
    WHERE model = c.model
)
```

The query above find the oldest cards of every model.

### Subqueries in FROM Clause

```
SELECT owner FROM
Cars c, (SELECT * FROM OwnerInsurance WHERE payment <
200) cheapInsOwn
WHERE c.owner = cheapInsOwn.owner;
```

The query above finds all cars with their respective insurance whose owner has a cheap insurance.

## Full Relation Operations

### Removing Duplicates

To remove duplicates use a the DISTINCT keyword as in SELECT DISTINCT .... But note that it is very expensive as all the relations will need to be sorted in order to detect and eliminate duplicates. Note that NULL will also be included at most once if it occurs in the result.

### Recursive Query

The basic structure is

```
WITH RECURSIVE recRel AS(
non_recursive-term
UNION [ALL]
recursive-term
) SELECT * FROM recRel;
```

The database executes the recursive-term as long as the resulting set keeps changing.

```
WITH RECURSIVE ancestor AS (
    SELECT * FROM people
UNION
    SELECT * from people as p
    WHERE p.father = ancestor.id OR p.mother =
        ancestor.id
)

SELECT * from ancestor;
```

### Aggregation Operators

SQL has five aggregation operators: SUM, AVG, MIN, MAX, COUNT. All except the COUNT are applied to scalar valued columns. The only excpetion being COUNT(\*) which counts the tuples in the relation constructed from the FROM and WHERE clause.

```
SELECT AVG(netWorth), COUNT(*)
FROM MovieExec;
```

### Grouping

The GROUP BY clause follows the WHERE clause and lists arguments that, if two tuples match on these arguments, will place them in the same group. The aggregation operators in the SELECT are applied on a per group basis. Only attributes mentioned in the GROUP BY clause may appear in the SELECT clause.

```
SELECT studioName, SUM(length)
FROM Movies
GROUP BY studioName;
```

### Grouping, Aggregation and Nulls

- The value NULL is ignored in aggregation. COUNT(\*) though counts number of tuples, but COUNT(A) counts number of attributes A and skips NULLS.
- NULL is treated as an ordinary value when forming groups.
- Aggregations over empty bag are NULL, except COUNT(\*) being 0.

### Having

Assume we would like to filter the groups we have aggregated: We can do this by adding a HAVING clause. e.g. we only want the names of sellers which have no sale below or equals 500:

```
SELECT COUNT(*), Name
FROM SellerSale
GROUP BY Name
HAVING MIN(Sale) > 500;
```

## Data Modifications

### Insertions

Basic syntax:

INSERT INTO  $R(A_1, \dots, A_n)$  VALUES  $(v_1, \dots v_n)$ ;

if they are already in standard order we can also write:

INSERT INTO R VALUES  $(v_1, \dots v_n)$ ;

or insert programmatically:

```
INSERT INTO Relation1
SELECT ....
```

We note that values are only inserted after the query has been fully computed preventing interference.

### Deletions

We cannot specify the exact tuple to be deleted, instead we must specify a set of tuples (which can contain 1 tuple) to be deleted.

```
DELETE FROM R WHERE <cond>
```

### Updates

Basic syntax:

UPDATE R SET <new vals> WHERE <cond>;

where the new vals are the name of the variable followed by an equals sign and the new assignment.

```
UPDATE Professor
SET name = 'Prof.┐' || name
WHERE title = 'Professor';
```

recall that the || operator concatenates strings.

## Transactions

A transaction is a set of SQL commands that get executed atomically. In PostgreSQL every statement in the shell is automatically wrapped in a transaction:



```
BEGIN;
UPDATE accounts SET balance = balance - 100.00
WHERE name = 'Alice';
SAVEPOINT my_savepoint;
UPDATE accounts SET balance = balance + 100.00
WHERE name = 'Bob';
-- oops ... forget that and use Wally's account
ROLLBACK TO my_savepoint;
UPDATE accounts SET balance = balance + 100.00
WHERE name = 'Wally';
COMMIT;
```

## Views

A view is a virtual relation which is built on a query run on a table.

```
CREATE VIEW <viewName> as <viewDefinition>;
```

where viewDefinition is a SQL query.  
we can delete a view with DROP VIEW <viewName>;

### Updating Views

We can only update views satisfying some constraints:

- only one base relation involved
- the view involves the key of the base relation
- view does not involve aggregates, group by or duplicate-elimination

## Indexes

To add an index:  
CREATE INDEX <indexName> ON Table1(field1);  
to remove an index:  
DROP INDEX <indexName>

### Multicolumn Indexes

CREATE INDEX <indexName> on Table1(field1,field2);  
Multicolumn indexes first index on field1 and then on field2.

## Isolation Level

With the comman SET TRANSACTION ISOLATION LEVEL <isolation level> we can set the isolation level for the current transaction.  
We can also have a new transaction with a specific isolation level by typing BEGIN TRANSACTION ISOLATION LEVEL <isolation level>;  
There are the following levels:

	Dirty Reads	Nonrepeatable Reads	Phantoms
READ UNCOMMITTED			
READ COMMITTED	No		
REPEATABLE READ	No	No	
SERIALIZABLE	No	No	No

- where:
- Dirty Reads: When we read data another transaction has modified and not yet committed.
  - Nonrepeatable Reads: When we read a record twice but the value was allowed to change in between.
  - Phantoms: When another transactions commits and adds or removes rows to a table we are currently reading.

## Part II SQL for Table Management

## Table Creation

The basic syntax is as follows:

```
CREATE TABLE Professor(  
    PersNr integer,  
    Name varchar(30),  
    Level character(2) default "AP",  
    PRIMARY KEY (PersNr)  
);
```

The default is a value that will be used if no Level is specified (otherwise we would have a NULL).

### Data Types

- int, bit integer, 0-1 bit
- numeric(precision) a number with fixed precision
- boolean
- timestamp, date, time
- varchar(n) variable length string with limit
- char(n) string which will be padded to size n
- text unlimited string (in theory)

## Relations

### Declaring Foreign Keys

We have two options:

```
CREATE TABLE Dog(  
    name VARCHAR(20),  
    owner INT REFERENCES Person(id)  
);
```

```
CREATE TABLE Dog(  
    name VARCHAR(20),  
    owner INT,  
    FOREIGN KEY (owner) REFERENCES Person(id)  
);
```

With the second version we can also define composite foreign keys by writing  
FOREIGN KEY (<attributes>) REFERENCES <table>(<attributes>)

### Referential Integrity

What should happen to a record (called child record in the following discussion) if a parent record (the one referred to in one of the childs foreign key fields) gets updated or deleted?  
There are multiple options: (usually the key a child refers to is a parents primary key)

- CASCADE
  - UPDATE if the parents primary key is updated the child’s foreign key will also be updated.
  - DELETE if parent gets deleted, child will be deleted too.
- RESTRICT
  - UPDATE attempt to update parent’s primary key will fail with an error
  - DELETE attempt to update parent’s primary key will fail with an error
- NO ACTION
  - UPDATE attempt to update parent’s primary key will fail with an error
  - DELETE attempt to update parent’s primary key will fail with an error
- SET DEFAULT
  - UPDATE when parent’s primary key gets updated child’s foreign key will be set to DEFAULT or NULL
  - DELETE when parent gets deleted child’s foreign key will be set to DEFAULT or NULL

**Difference between NO ACTION and RESTRICT?** With RESTRICT the constraint will be checked and immediately fails whereas with NO ACTION the database tries to update/delete but then fails on committing (failure of deferred check). However some databases (e.g. MySQL) do not differentiate between the two options.  
Usually NO ACTION is the **default** if no specific constraint has been chosen.  
The constraints can be specified during table creation as follows:

```
create table Lecture (  
    ...,  
    PersNr integer references Professor  
on delete {cascade | restrict | no action |  
    set null |set default}  
on update {cascade | restrict | no action |  
    set null |set default}  
);
```

## Enforcing 1:1 Relations

This is best done using a third table:

```
CREATE TABLE Employee(  
    emp_id VARCHAR(20) NOT NULL UNIQUE  
);  
CREATE TABLE Office(  
    room VARCHAR(20) NOT NULL UNIQUE  
);  
  
CREATE TABLE HasOffice(  
    emp_id VARCHAR(20) NOT NULL UNIQUE REFERENCES  
        Employee (emp_id),  
    room VARCHAR(20) NOT NULL UNIQUE REFERENCES Office  
        (room)  
);
```

### Constraints

#### NOT NULL

Attribute cannot be NULL

```
..  
sName text NOT NULL,  
...
```

#### PRIMARY KEY

Only one primary key constraint can be specified per relation. Two options, either a single attribute as primary key like

```
...  
sID int PRIMARY KEY,  
...
```

or a combination of attributes:

```
CREATE TABLE Student  
(  
    sID int,  
    sName text,  
    GPA real  
    PRIMARY KEY (sID, sName)  
);
```

where the combination must be unique and all attributes need to be non null. **UNIQUE**

Usage like primary key: A value must be unique, but different from PRIMARY KEY the value can be NULL and we can have multiple UNIQUE constraints per relation.

#### CHECK

Can happen at tuple or attribute level which will get checked at update of tuple or single attribute respectively.

```
CREATE TABLE Employee  
(  
    salary real,  
    bonus real,
```

```
CHECK(salary + bonus > 10000)  
);
```

```
CREATE TABLE Employee  
(  
    salary real CHECK(salary > 5000)  
);
```

### Assertions & Triggers

#### Assertion

A logical expression that must be true at all times:

```
CREATE  ASSERTION <assertName> CHECK (<condition>)
```

#### Trigger

A trigger is an SQL statement that gets executed when a certain predefined action is performed. It is best to just look at an example or read the docs if a trigger is needed:

```
CREATE TRIGGER CeoCantGetPayCut  
AFTER UPDATE OF (salary) ON Ceo  
REFERENCING  
    OLD ROW AS oldRow,  
    NEW ROW AS newRow,  
FOR EACH ROW  
WHEN(oldRow.salary > newRow.salary)  
    UPDATE Ceo  
    SET salary = oldRow.salary  
    WHERE id = oldRow.id;
```